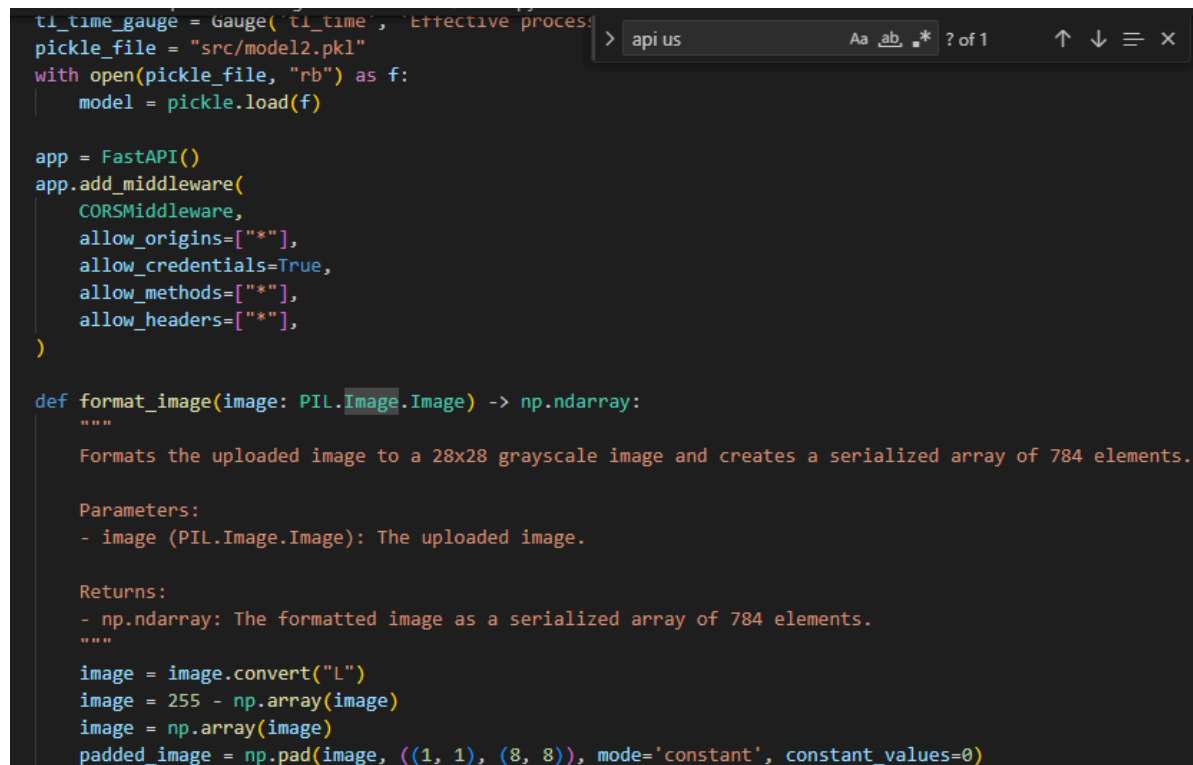


# Report NA19B056 Assignment 7

## Prometheus and Grafana Monitoring for FastAPI Application

Screenshots are within the report

This report outlines the setup and usage of Prometheus and Grafana for monitoring a FastAPI application. The objective is to monitor critical metrics such as API network I/O, CPU utilization, and memory usage.



```
ti_time_gauge = Gauge('ti_time', 'Effective process')
pickle_file = "src/model2.pkl"
with open(pickle_file, "rb") as f:
    model = pickle.load(f)

app = FastAPI()
app.add_middleware(
    CORSMiddleware,
    allow_origins=["*"],
    allow_credentials=True,
    allow_methods=["*"],
    allow_headers=["*"],
)

def format_image(image: PIL.Image.Image) -> np.ndarray:
    """
    Formats the uploaded image to a 28x28 grayscale image and creates a serialized array of 784 elements.

    Parameters:
    - image (PIL.Image.Image): The uploaded image.

    Returns:
    - np.ndarray: The formatted image as a serialized array of 784 elements.
    """
    image = image.convert("L")
    image = 255 - np.array(image)
    image = np.array(image)
    padded_image = np.pad(image, ((1, 1), (8, 8)), mode='constant', constant_values=0)
```

### Setting Up Prometheus

Prometheus is an open-source monitoring and alerting toolkit. It collects and stores metrics as time series data, providing a powerful query language (PromQL) to analyze this data.

#### Installing Prometheus

1. Run Prometheus with the configuration file:

```
prometheus --config.file="/etc/prometheus/prometheus.yml"
```

## Configuring Prometheus

### 1. Modify the Prometheus configuration file:

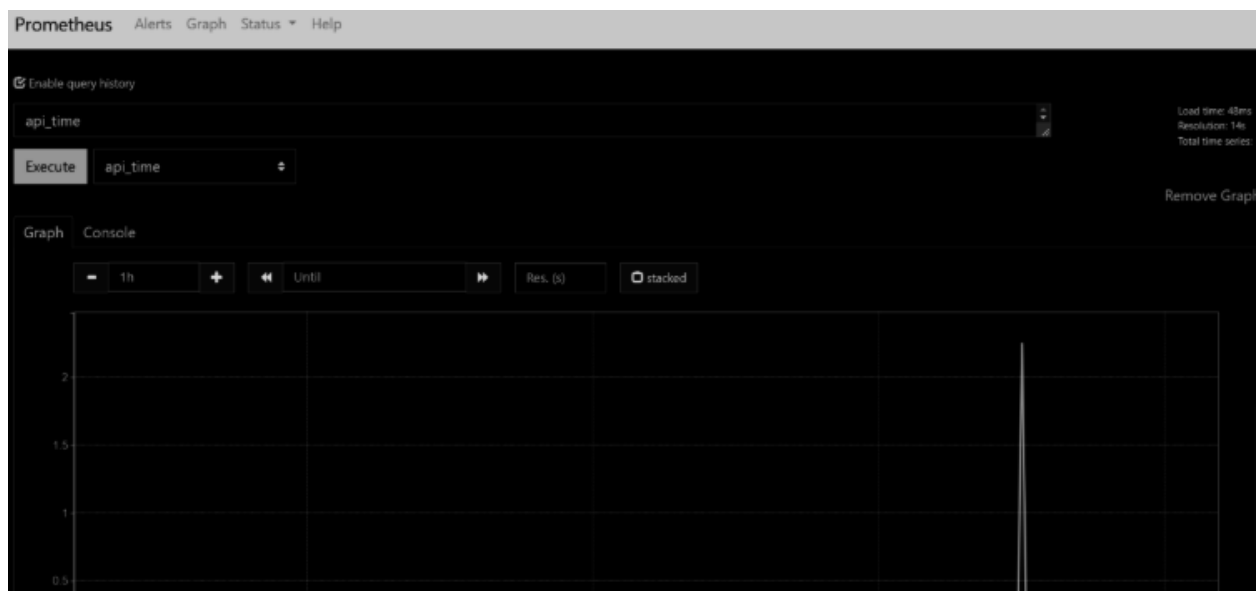
```
prometheus --config.file="/etc/prometheus/prometheus.yml"
```

Added the following scrape configuration to monitor the FastAPI application:

```
scrape_configs:
  - job_name: 'fastapi'
    static_configs:
      - targets: ['localhost:8000']
```

### 2. Start Prometheus:


```
prometheus --config.file=/etc/prometheus/prometheus.yml
```



## Setting Up Grafana

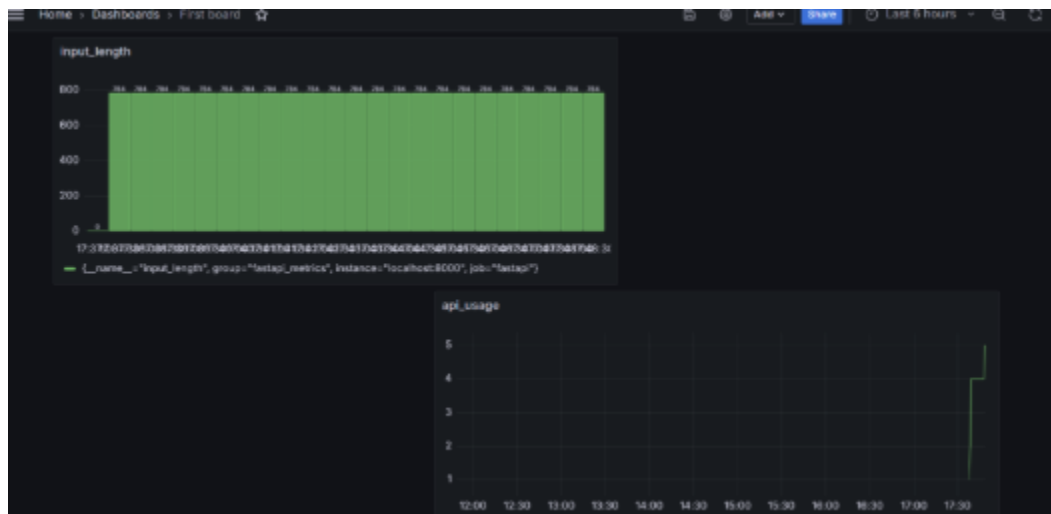
Grafana is an open-source platform for monitoring and observability. It allows you to query, visualize, alert on, and explore metrics from multiple data sources, including Prometheus.


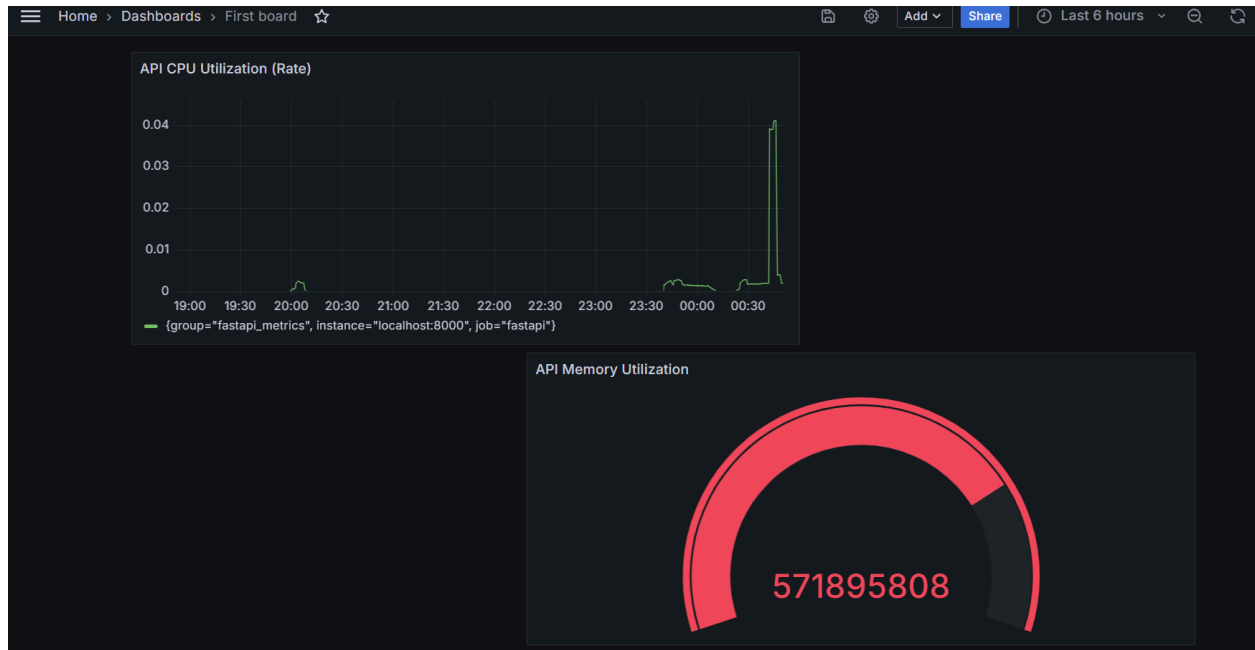
# Grafana



<tc.onyemaobi@gmail.com>

Source: Tech and Beyond With Moss





## Installing Grafana

### 1. Download and install Grafana:

```
wget https://dl.grafana.com/oss/release/grafana-7.3.6.linux-amd64.tar.gz
tar -zxvf grafana-7.3.6.linux-amd64.tar.gz
cd grafana-7.3.6
```

### 2. Start Grafana:

```
./bin/grafana-server
```

## Configuring Grafana

1. **Access Grafana:** Navigate to local host in your web browser.
2. **Login:** Default credentials are `admin / admin`.
3. **Add Prometheus as a data source:**
  - Go to Configuration > Data Sources > Add data source.
  - Select Prometheus and configure it with the local host URL.
4. **Create a Dashboard:** and add metrics you want to monitor.

## Dockerizing the FastAPI Application

Docker allows you to package your application and its dependencies into a container, ensuring consistency across different environments.

1. **Create a Dockerfile:**

```
FROM tiangolo/uvicorn-gunicorn-fastapi:python3.8  
  
COPY ./app /app  
  
RUN pip install --no-cache-dir -r /app/requirements.txt
```

2. **Build the Docker image:**

```
docker build -t fastapi_app
```

3. **Run the Docker container:**

```
docker run -d -p 8000:8000 fastapi_app
```

## Monitoring Metrics with Prometheus

Prometheus scrapes metrics from the FastAPI application at regular intervals. The following metrics are collected and monitored:

## API Network I/O Bytes

### 1. Rate of API network I/O bytes:

`rate(http_request_size_bytes_created[5m]) - rate(http_response_size_bytes_created[5m])`

### 2. Total API network I/O bytes:

```
http_request_size_bytes_total - http_response_size_bytes_total
```

## API CPU Utilization

### 1. Rate of CPU utilization:

```
rate(process_cpu_seconds_total{job="fastapi", instance="localhost:8000"}[5m])
```

## API Memory Utilization

### 1. Memory utilization:

```
process_resident_memory_bytes{instance="localhost:8000", job="fastapi"}
```

## Additional Metrics

### 1. API Time: Monitor the time taken by the API to respond.

```
api_time
```

### 2. Task Length Time: Monitor the time taken to complete specific tasks.

```
tl_time
```

### 3. Input Length: Monitor the length of input data.

```
input_length
```

### 4. API Usage: Sum of API usage by client IP.

```
sum(api_usage_total) by (client_ip)
```

Github Link: <https://github.com/Rinkle-S/CS5830/tree/main/A07>